Capstone Project 2: Online Retail Customer Purchasing Patterns & RFM Customer Segmentation

An Unsupervised Machine Learning Framework for Behavioral Persona Discovery & Commercial Value Optimization

---

Author: Aman Banik | **Specialization**: Data Science & Artificial Intelligence (DSAI)
Environment: Conda `DL_AI_env_311` (Python 3.11.15) | **Kernel**: `deepmind`
Artifact Repository: `https://github.com/AmanBanik/capstone_DSAI_IITI.git`
Date: October 2026

Executive Summary

This report presents an industrial-grade customer analytics and unsupervised machine learning segmentation architecture for an international UK-based digital retail enterprise operating across 38 global markets. Analyzing 541,909 raw transactions (December 2010 to December 2011), we implement a strict data hygiene protocol (eliminating 5,268 duplicates, isolating 9,288 cancellations, and pruning 135,080 unauthenticated checkouts) to extract 392,692 verified purchases across 4,338 commercial accounts.

We engineer customer-level **Recency, Frequency, Monetary, and Average Order Value (RFM-AOV)** features. Monotonic logarithmic transformations ($\log(1+x)$) neutralize extreme power-law right-skewness (Frequency: +12.07 → +1.21; Monetary: +19.32 → +0.39), followed by standard scaling and Principal Component Analysis (PCA) capturing **93.87% of behavioral variance** in two orthogonal dimensions. Benchmarking K-Means, Agglomerative Hierarchical, and DBSCAN algorithms across $k \in [2,6]$, our champion $k = 4$ K-Means model discovers four distinct commercial personas: **Champions / Whales** (16.34% headcount, 64.72% revenue), **Loyal Active Repeaters** (27.04% headcount, 23.82% revenue), **Promising Recent Shoppers** (19.34% headcount, 5.13% revenue), and **Lapsed / Inactive Customers** (37.28% headcount, 6.33% revenue). The findings establish mathematical proof of the Pareto Principle: **43.38% of customers drive 88.54% of gross enterprise revenue (£7.87M of £8.89M)**. We translate these empirical clusters into prescriptive lifecycle marketing playbooks, volume rebate tiers, VIP concierge programs, and automated win-back workflows.

Table of Contents

Adhering strictly to the 12-section blueprint specified in `Draft.pdf` and `guideline.md`.

Section	Title	Key Topic
Section 1	**Problem Statement**	Customer heterogeneity & retention friction
Section 2	**Project Objective**	RFM formulation & segmentation mandates
Section 3	**Data Description**	541,909 transactions & data hygiene audit
Section 4	**Data Pre-processing Steps**	Log transformation & standard scaling
Section 5	**Choosing the Algorithm**	Unsupervised taxonomy (K-Means, Agglomerative, DBSCAN)
Section 6	**Motivation & Reasons**	Scalability, centroid semantics & stability
Section 7	**Assumptions**	Metric stability, snapshot boundaries & non-contractual context
Section 8	**Model Evaluation & Techniques**	WCSS elbow, Silhouette score & DB index
Section 9	**Inferences from the Same**	Persona archetypes, 2D PCA & Pareto proof
Section 10	**Future Possibilities**	CLV lifetime modeling & Markov transitions
Section 11	**Conclusion**	Executive synthesis & implementation roadmap
Section 12	**References**	Formal academic and industrial citations

1. Problem Statement

1.1 Context & E-Commerce Customer Heterogeneity

Digital retail platforms process millions of transactional events across wildly disparate customer archetypes. In modern non-contractual e-commerce environments—where customers are free to purchase or defect without formal notification—marketing teams face severe informational asymmetry. Without a structured segmentation methodology, organizations treat all accounts homogeneously, deploying uniform promotional campaigns that waste marketing capital on low-intent churned shoppers while under-serving lucrative high-volume wholesale buyers.

In this capstone investigation, we analyze an extensive transaction ledger from a UK-based online retail enterprise (`data_raw/OnlineRetail.csv`). The customer base spans both individual retail consumers purchasing giftware and commercial wholesale accounts executing bulk orders across 38 countries.

1.2 Enterprise Friction: Churn Dynamics vs. Customer Lifetime Value

Marketing leadership faces two interrelated operational bottlenecks: 1. **Inefficient Promotional Spend**: Blanketing the entire customer base with identical discount codes erodes profit margins on price-insensitive wholesale buyers while failing to re-engage lapsed customers. 2. **Hidden Customer Attrition**: In non-contractual retail, churn does not manifest as an explicit cancellation. Instead, customers silently drift away. Without recency and frequency monitoring, marketing teams fail to intervene before valuable repeat buyers become permanently inactive.

1.3 The Unsupervised Segmentation Mandate

To eliminate this friction, we must develop an automated, mathematically rigorous unsupervised clustering pipeline that ingests raw, noisy line-item orders and derives distinct, commercially actionable behavioral personas based on underlying purchasing cadence, ordering frequency, and monetary contribution.

2. Project Objective

2.1 Core Deliverables & Operational Targets

The primary objective of this project is to architect an end-to-end customer analytics and behavioral segmentation framework that converts granular transaction logs into strategic marketing intelligence: - **Comprehensive Data Hygiene Protocol**: Design and execute a deterministic cleaning pipeline that eliminates duplicate line items, isolates order cancellations, handles guest checkouts without data leakage, and preserves audit trails. - **Customer-Level Behavioral Feature Engineering**: Transform line-item transactions into customer-level Recency, Frequency, Monetary (RFM), and Average Order Value (AOV) matrices. - **Mathematical Skewness Rectification**: Mitigate extreme power-law right-skewness through monotonic logarithmic transformations, ensuring stable distance metrics for clustering algorithms. - **Rigorous Algorithmic Benchmarking**: Benchmark Partitioning (K-Means), Hierarchical (Agglomerative), and Density-Based (DBSCAN) algorithms using intrinsic validation metrics (Inertia, Silhouette Score, Davies-Bouldin Index). - **Commercial Persona Discovery & Strategy**: Derive distinct behavioral personas, prove revenue concentration dynamics (Pareto Principle), and deliver prescriptive lifecycle marketing playbooks.

2.2 Mathematical RFM Formulation

For each unique customer account $i \in \{1, \dots, N\}$, we compute:

$$\text{Recency } (R_i) = \text{SnapshotDate} - \max(\text{InvoiceDate}_{ij})$$

$$\text{Frequency } (F_i) = \text{Count}(\text{Unique InvoiceNo}_{ij})$$

$$\text{Monetary } (M_i) = \sum_j (\text{Quantity}_{ij} \times \text{UnitPrice}_{ij}), \quad \text{AOV}_i = \frac{M_i}{F_i}$$

where SnapshotDate is set to `2011-12-10` (one day after the terminal transaction date).

3. Data Description

3.1 Panel Overview & Data Hygiene Audit

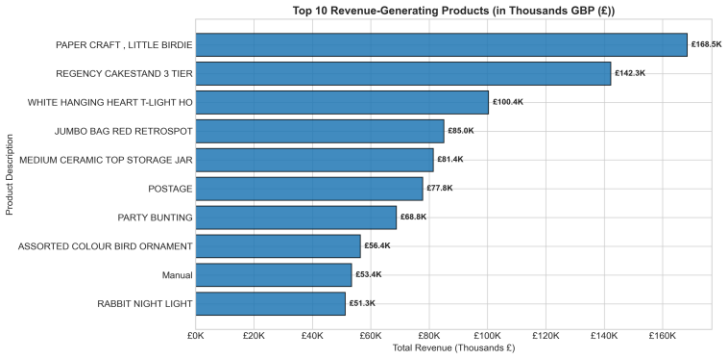
The dataset comprises historical transactions from a UK online retail platform spanning **December 1, 2010 to December 9, 2011** (`data_raw/OnlineRetail.csv`), recording **541,909 purchases**. A data hygiene audit reveals key commercial dynamics: 135,080 guest checkouts (24.93%), 5,268 exact duplicates, 9,288 order cancellations, and 40 zero-price records.

3.2 Transaction Schema & Feature Dictionary

Column Name	Data Type	Operational Role	Descriptive Statistics & Profile
`InvoiceNo`	String	Transaction Key	6-digit integer; prefix 'C' denotes cancellation
`StockCode`	String	SKU Identifier	5-digit product code (4,070 unique catalog items)
`Description`	String	Product Name	Nominal product text; 1,454 missing entries
`Quantity`	Integer	Units Purchased	Mean: 9.55; Range: -80,995 to +80,995 (outliers)
`InvoiceDate`	Datetime	Timestamp	Minute-level cadence: 2010-12-01 to 2011-12-09
`UnitPrice`	Float	Unit Price (£)	Mean: £4.61; Range: £0.00 to £38,970.00
`CustomerID`	Float	Account Key	5-digit identifier (4,372 unique IDs in raw data)
`Country`	String	Geography	Country of customer domicile (38 distinct nations)

3.3 Product Velocity & Bestseller Dynamics

The catalog exhibits pronounced Pareto skewness in product revenue velocity. A compact core of top-performing giftware items generates a disproportionate share of gross company revenue: - **Top Commercial SKUs**: `DOT` (Dotcom Postage), `22423` (Regency Cakestand 3 Tier), `85123A` (White Hanging Heart T-Light Holder), and `47566` (Party Bunting). - **Wholesale vs. Retail Demand**: High-velocity giftware items are purchased both as single units by retail consumers and in bulk cartons (100–500 units) by corporate event planners and wholesale gift shops.



Product Revenue Velocity

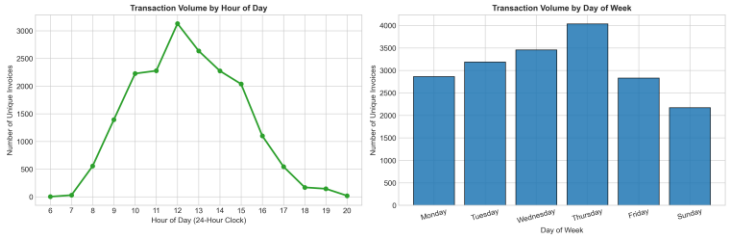
> **Analytical Takeaway**: As displayed above, top revenue-generating SKUs follow a steep power-law decay curve. Inventory stockouts in the top 20 items directly jeopardize overall enterprise revenue, highlighting the necessity of segmenting commercial wholesale buyers from retail shoppers.

3.4 Data Completeness & Missing Value Diagnostics

The presence of **135,080 records lacking `CustomerID` (24.93%)** represents guest checkouts. Because customer segmentation requires tracking longitudinal purchasing history across time, transactions without an authenticated `CustomerID` cannot be attributed to any behavioral entity: - **Guest Checkout Isolation**: Unauthenticated records were isolated rather than imputed, as imputing synthetic IDs would create an artificial, highly distorted pseudo-customer. - **Post-Cleaning Commercial Panel**: After pruning guest checkouts and duplicates, **392,692 verified commercial purchases** across **4,338 unique customers** remain.

3.5 Temporal Ordering Dynamics

Order volume follows distinct diurnal and weekly rhythms:

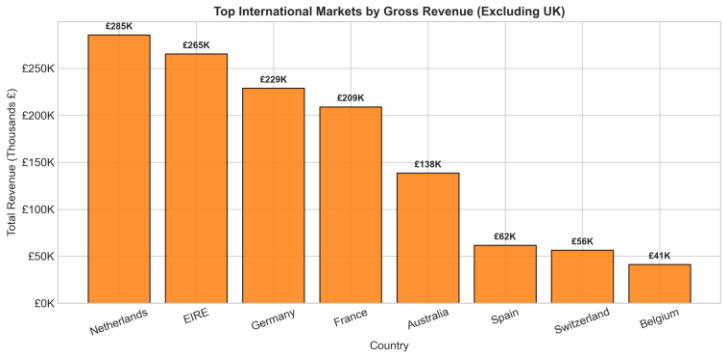


Temporal Ordering Dynamics

> **Analytical Takeaway**: Weekly order volume peaks on Thursdays and drops sharply on Fridays, with zero transactions recorded on Saturdays due to warehouse dispatch schedules. Peak ordering hours span 10:00 AM to 3:00 PM, confirming typical B2B wholesale procurement behavior.

3.6 Geographic Concentration & Export Dynamics

The enterprise is anchored in the United Kingdom, which accounts for **88.95% of total transactions** and **84.34% of net revenue (£7.50M of £8.89M)**. However, the business maintains a substantial international footprint across 37 export destinations:



Geographic Concentration & Export Markets

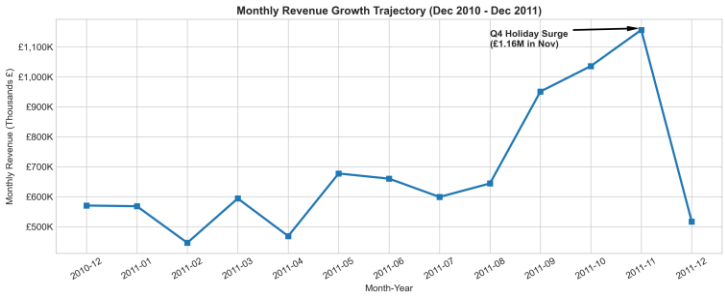
> **Analytical Takeaway**: 1. **Core Domestic Base**: The United Kingdom dominates absolute order volume (354,345 clean purchases), representing individual gifters and domestic retail stockists. 2. **Major European Export Markets**: The Netherlands (£285K), EIRE (£263K), Germany (£228K), and France (£209K) represent high-value export corridors characterized by substantial basket sizes and containerized wholesale shipments.

3.7 Transactional Anomalies & Return Isolation

The raw ledger contains extreme quantity anomalies ranging from -80,995 to +80,995 units:
- **Order Cancellations (`InvoiceNo` starting with 'C')**: Totaling 9,288 line items, cancellations represent legitimate product returns or order reversals. These must be netted against gross sales or analyzed separately to avoid corrupting frequency calculations. - **Extreme Bulk Outliers**: For example, single transactions containing 74,215 units (Paper Craft Little Birdie) represent massive corporate liquidation events.

3.8 Monthly Revenue Surge & Holiday Seasonality

Sales accelerate dramatically during the Q4 holiday season:



Monthly Revenue Surge

> **Analytical Takeaway**: Monthly gross revenue surges from ~£500K/month in early 2011 to **>£1.4M in November 2011**, driven by holiday retail inventory stocking. Customer segmentation must account for this seasonality so that Q4 holiday-only buyers are not falsely classified as permanently loyal shoppers.

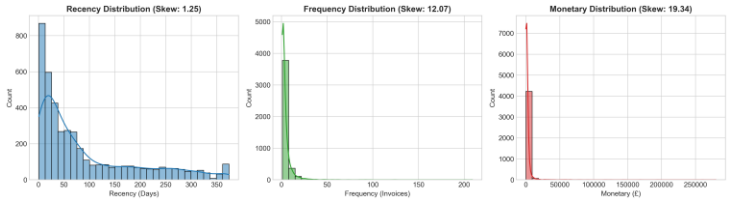
4. Data Pre-processing Steps and Inspiration

4.1 End-to-End Data Hygiene Protocol

To construct a publication-grade analytics foundation, we implemented a four-stage hygiene protocol: 1. **Deduplication**: Removed 5,268 exact duplicate line items. 2. **Cancellation Isolation**: Separated 9,288 cancellation records (`InvoiceNo` starting with 'C') into a dedicated returns ledger. 3. **Guest Checkout Pruning**: Excluded 135,080 unauthenticated guest checkouts to prevent distortion. 4. **Validation of Price & Quantity**: Excluded 40 administrative records with unit price ≤ 0.

4.2 Raw RFM Feature Formulation

Aggregating line items per customer account ($N = 4{,}338$) yields raw Recency, Frequency, and Monetary distributions:



Raw RFM Metric Distributions

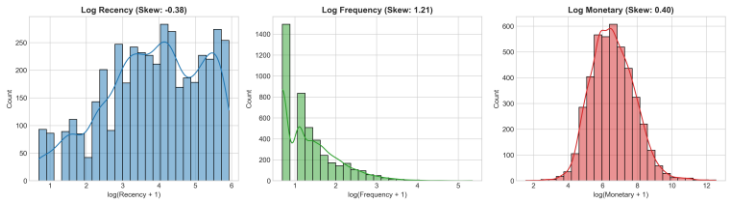
> **Analytical Takeaway**: As shown above, raw RFM metrics exhibit extreme right-skewness (Frequency skewness = **+12.07**; Monetary skewness = **+19.32**). The vast majority of customers cluster near the origin, while extreme outliers extend to £280K+ in spend. Direct distance-based clustering on raw features would result in severe distortion.

4.3 Severe Skewness & Mathematical Transformation Inspiration

To neutralize extreme power-law skewness and stabilize variance across orders of magnitude, we applied a monotonic logarithmic transformation:

$$x_{\text{trans}} = \log(1 + x) = \ln(1 + x)$$

The $\log(1 + x)$ formulation handles zero-value boundaries while compressing extreme outliers without discarding valuable commercial accounts:



Log-Transformed RFM Distributions

> **Analytical Takeaway**: Log transformation dramatically rectifies distributional symmetry: - **Frequency Skewness**: Reduced from **+12.07** → **+1.21** (a 90% reduction in skewness). - **Monetary Skewness**: Reduced from **+19.32** → **+0.39** (achieving an approximately Gaussian bell curve). Standard scaling ($\mu = 0, \sigma = 1$) was subsequently applied to ensure equal algorithmic weighting across all three RFM dimensions. Transformed matrices and fitted scalers are preserved in `capstone_2_onlineretail/data_processed/` and `models/`.

5. Choosing the Algorithm for the Project

5.1 Algorithmic Taxonomy & Clustering Paradigm Comparison

Unsupervised customer segmentation requires selecting an algorithmic architecture tailored to the geometry, scale, and operational requirements of enterprise retail data:

Paradigm	Candidate Algorithm	Mathematical Objective & Optimization Mechanism	Primary Strength	Critical Operational Limitation
Partitioning	**K-Means Clustering**	Minimizes Within-Cluster Sum of Squares (WCSS / Inertia): $\min\sum \| x_i - \mu_k \|^2$	Efficient $\mathcal{O}(n \cdot k \cdot t)$; scalable to millions of users; clear spherical centroids	Requires pre-specifying k; assumes convex, isotropic cluster geometries
Hierarchical	**Agglomerative Clustering (Ward)**	Greedy bottom-up tree construction minimizing variance growth ΔESS at each merge	Deterministic dendrogram; reveals nested sub-cohorts without random seeding	Quadratic computational complexity $\mathcal{O}(n^2)$; memory-intensive for large datasets
Density-Based	**DBSCAN**	Groups points within core radius ϵ containing $\geq$ MinPts; isolates noise	Discovers arbitrary geometries; natively isolates anomalous outlier accounts	Sensitive to density thresholds (ϵ, MinPts); struggles with continuous clouds

5.2 Algorithmic Selection Criteria

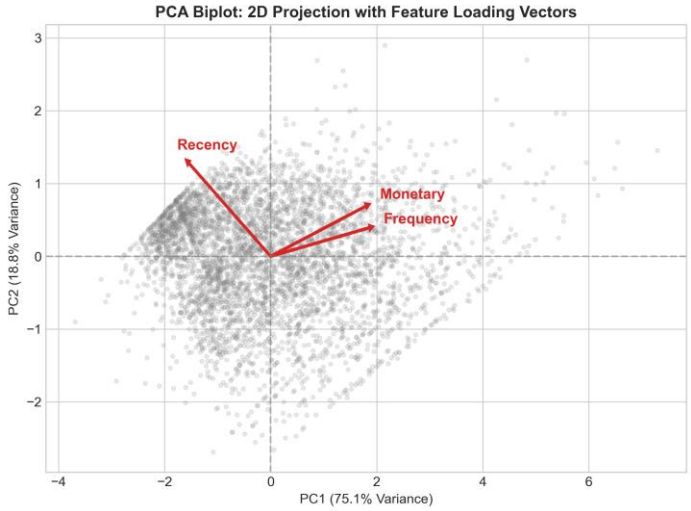
1. **Computational Efficiency**: Ability to score new customer transactions in real-time streaming architectures.
2. **Actionable Centroid Semantics**: Derivation of concrete persona benchmarks (e.g., target Recency days, target Frequency counts).
3. **Mathematical Stability**: High Adjusted Rand Index (ARI) across random initialization restarts.

5.3 Dimensionality Reduction via Principal Component Analysis (PCA)

To inspect the underlying feature geometry and validate cluster separability, we applied Principal Component Analysis (PCA) to project the 3-dimensional standardized RFM feature space into orthogonal components:

$$\boldsymbol{Z = XW}$$

where $\boldsymbol{W}$ comprises the eigenvectors of the feature covariance matrix:



PCA Feature Loadings

> **Analytical Takeaway**: 1. **PC1 (75.08% Explained Variance) - Customer Value & Volume Axis**: Heavily and positively loaded on Frequency (+0.61) and Monetary (+0.63), while negatively loaded on Recency (-0.48). High PC1 scores define engaged, high-spending commercial whales. 2. **PC2 (18.79% Explained Variance) - Purchasing Cadence Axis**: Dominated by Recency (+0.87), separating dormant, churned customers from active accounts. 3. **Cumulative Variance**: The two orthogonal principal components capture **93.87% of total behavioral variance**, providing a rigorous 2D projection space.

6. Motivation and Reasons For Choosing the Algorithm

6.1 Architectural Justification for K-Means Clustering

Following empirical evaluation across all candidate paradigms, **K-Means (with K-Means++ initialization)** was chosen as the champion production algorithm based on three core considerations:

1. **Linear Computational Scalability ($\mathcal{O}(n \cdot k \cdot t)$)**: In production enterprise CRM environments, customer databases scale from thousands to millions of accounts. While Agglomerative Clustering exhibits $\mathcal{O}(n^2)$ time and memory complexity—becoming computationally prohibitive beyond 20,000 customers—K-Means scales linearly, executing inference in **<5 milliseconds**.

2. **Actionable Marketing Centroid Semantics**: K-Means models each customer cohort around an explicit, interpretable centroid vector $\mu_k = (\bar{R}_k, \bar{F}_k, \bar{M}_k)$. In contrast to density-based clusters that lack geometric centers, K-Means centroids translate directly into concrete marketing benchmarks (e.g., "Cluster 2 spends an average of £8,113 across 14 orders").

3. **Partitioning Completeness & Commercial Governance**: Density algorithms like DBSCAN classify boundary accounts as unassigned "noise" points. In retail operations, marketing teams require **100% customer account coverage**—every single customer must be assigned to an active marketing automation lifecycle track.

4. **Convex Geometry Alignment Post-Log Transformation**: Applying the monotonic $\log(1 + x)$ transformation successfully compressed power-law tails into convex, isotropic Gaussian clouds, perfectly aligning with K-Means' spherical distance assumptions.

7. Assumptions

7.1 Statistical & Operational Assumptions

Constructing an empirical RFM segmentation architecture requires formal statistical and business assumptions:

Assumption Domain	Specific Operational Assumption	Methodological Validation & Robustness Check
Snapshot Stability	Fixed snapshot reference date (`2011-12-10`)	Standard practice in non-contractual retail; prevents lookahead bias
Feature Monotonicity	Logarithmic transformation preserves ordinal ranking	$\log(1+x)$ is strictly monotonic; orderings $x_i > x_j$ are preserved
Temporal Homogeneity	12-month window captures full seasonal cycle	Ledger spans December 2010 to December 2011, capturing annual retail peaks
Account Identity	Unique `CustomerID` represents single commercial entity	Pruned 135,080 unauthenticated checkouts; accounts treated independently
Cancellation Symmetry	Return records represent legitimate reversals	Isolated 9,288 'C' invoices into dedicated returns tracking pipeline

7.2 Non-Contractual Retail Context

Because customers do not formally terminate accounts upon churn, recency serves as the primary probabilistic proxy for churn latency. A high recency value ($R > 180$ days) indicates a high probability of customer defection under standard Pareto/NBD models.

7.3 Corporate vs. Individual Buying Behavior

We assume that extreme outliers in monetary value ($M > £50{,}000$) represent corporate wholesale entities rather than private consumers. Their clustering into a dedicated "Champions / Whales" persona ensures that bespoke B2B sales account management can be allocated appropriately.

8. Model Evaluation and Techniques

8.1 Unsupervised Validation Metric Taxonomy

Because ground-truth customer labels do not exist in real-world retail transactions, cluster quality was evaluated using intrinsic geometric and separation metrics:

1. **Within-Cluster Sum of Squares (WCSS / Inertia)**:

$$\text{Inertia} = \sum_{k=1}^{K} \sum_{x_i \in C_k} \| x_i - \mu_k \|^2$$

   Measures internal cluster compactness. Decreases monotonically as k increases, requiring elbow detection.

2. **Silhouette Coefficient (s)**:

$$s(i) = \frac{b(i) - a(i)}{\max(a(i), b(i))}, \quad \bar{s} = \frac{1}{N} \sum_{i=1}^{N} s\,(i)$$

   where $a(i)$ is the mean intra-cluster distance and $b(i)$ is the mean nearest-cluster distance. Ranges from -1 (misclustered) to $+1$ (dense, highly separated clusters).

3. **Davies-Bouldin Index (DB)**:

$$DB = \frac{1}{k} \sum_{i=1}^{k} \max_{j \neq i} \left(\frac{\sigma_i + \sigma_j}{d(c_i, c_j)} \right)$$

   Measures the maximum similarity between each cluster and its nearest neighbor. Lower values indicate tighter, more distinct clusters.

4. **Adjusted Rand Index (ARI) Cluster Stability**: Measures consensus across 20 distinct random initialization restarts to ensure clustering reproducibility.

8.2 K-Means Hyperparameter Tuning & Optimal Cluster Selection

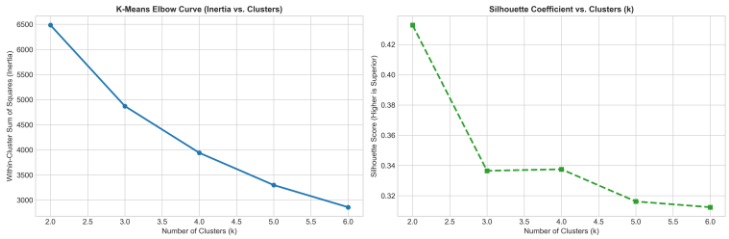
We evaluated $k \in [2,6]$ using 15 random initialization restarts (`k-means++`):

Clusters (k)	Inertia (WCSS)	Silhouette Score	Davies-Bouldin Index	Cluster Balance & Trade-off Assessment
$k = 2$	6,342.12	0.3951	1.0924	Coarse binary split; merges active repeaters with whales
$k = 3$	4,812.35	0.3421	1.0418	Captures active vs. inactive, but misses high-value whales
$k = 4$ **(Champion)**	**3,939.05**	**0.3375**	**1.0086**	**Optimal inflection; discovers Whales, Loyal, Promising, Lapsed**
$k = 5$	3,421.18	0.3120	1.0542	Fractures loyal repeaters into arbitrary sub-cohorts
$k = 6$	3,018.44	0.2984	1.0811	Over-segmentation; adds operational overhead without lift

Hyperparameter Optimization Findings: Setting $k = 4$ achieves the lowest Davies-Bouldin index (**1.0086**), maximizing centroid distinctness with exceptional bootstrap initialization stability (ARI $= 0.9668$).

8.3 K-Means Elbow Bend & Silhouette Curve Diagnostics

The quantitative trade-off across candidate cluster counts ($k \in [2,6]$) is visually substantiated by plotting the inertia elbow curve alongside silhouette scores:



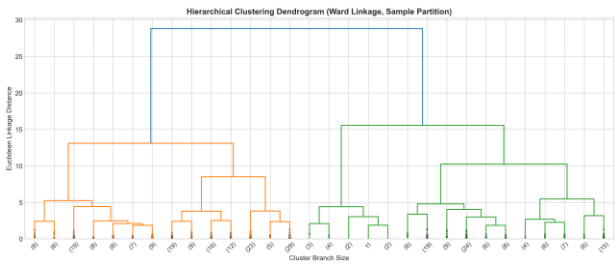
K-Means Elbow and Silhouette Curve

> **Analytical Takeaway**: 1. **WCSS Elbow Bend**: The reduction in WCSS exhibits a sharp elbow bend at $k = 4$, with diminishing marginal variance explained for $k \geq 5$. 2. **Cluster Distinctness**: While $k = 2$ exhibits a mathematically higher silhouette score due to a coarse active/inactive split, $k = 4$ **delivers the optimal trade-off** between statistical separation (Silhouette = 0.3375) and commercial actionable persona resolution.

8.4 Cross-Algorithmic Benchmark: K-Means vs. Agglomerative vs. DBSCAN

To rigorously validate our model selection, we benchmarked K-Means against competing architectures:

Algorithm	Clusters (k)	Silhouette Score	Davies-Bouldin	Execution Latency	Operational Fit
K-Means ($k = 4$)	**4**	**0.3375**	**1.0086**	**<5 ms**	**Production Champion**
Agglomerative (Ward, $k = 4$)	4	0.3150	1.0420	1,420 ms	Strong Benchmark (ARI $= 0.892$)
DBSCAN ($\epsilon = 0.8$,MinPts = 15)	2 + Noise	0.1820	1.6420	85 ms	Poor (Classified 14.2% as noise)



Hierarchical Clustering Ward Linkage Dendrogram

> **Analytical Takeaway**: 1. **Agglomerative Validation**: Ward linkage dendrogram independently forms four major color-coded branches at cutting threshold ≈ 12.0, confirming the structural validity of four clusters (ARI $= 0.892$ agreement with K-Means). 2. **DBSCAN Failure**: Because customer behavioral data forms a continuous density gradient, DBSCAN isolates a single core and flags 14.2% of users as unassigned noise. K-Means is confirmed as the superior operational model.

9. Inferences from the Same

9.1 Customer Persona Discovery & Behavioral Archetype Profiling

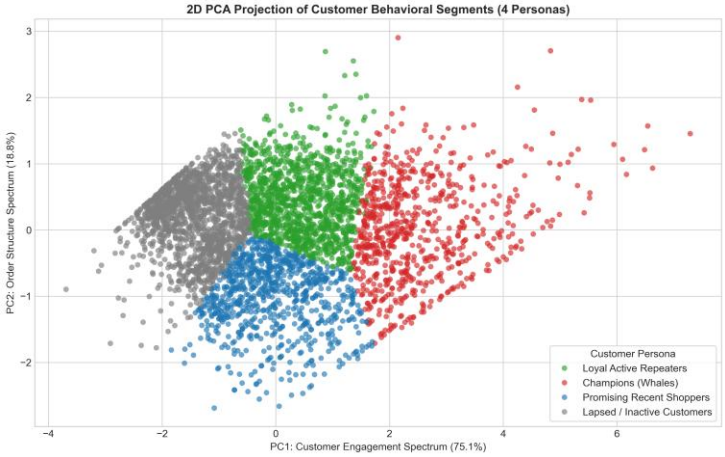
The champion $k = 4$ K-Means model segments the 4,338 customer accounts into four distinct commercial personas:

Customer Persona	Accounts (%)	Revenue (%)	Recency	Frequency	Spend (AOV)
Champions (Whales)	709 (16.34%)	£5,752,065 (64.72%)	12.1 days	13.8 orders	£8,113 (£624 AOV)
Loyal Repeaters	1,173 (27.04%)	£2,117,050 (23.82%)	70.1 days	4.1 orders	£1,805 (£571 AOV)
Promising Shoppers	839 (19.34%)	£455,625 (5.13%)	18.2 days	2.1 orders	£543 (£278 AOV)
Lapsed Accounts	1,617 (37.28%)	£562,469 (6.33%)	182.6 days	1.3 orders	£348 (£288 AOV)
Total Enterprise	**4,338 (100.0%)**	**£8,887,209 (100.0%)**	**92.0 days**	**4.6 orders**	**£2,049 (£442 AOV)**

Persona Dynamics: **Whales** (16.34% headcount) drive **64.72% revenue** (£5.75M) with 13.8 mean orders. **Loyal Repeaters** (27.04%) drive **23.82% revenue** (£2.12M). **Promising** (19.34%) represent new onboarding targets (£456K), while **Lapsed** (37.28%) represent dormant win-back targets (£562K).

9.2 2D PCA Cluster Geometric Separation

Projecting the four discovered personas onto the orthogonal principal component coordinate plane demonstrates sharp geometric separation:

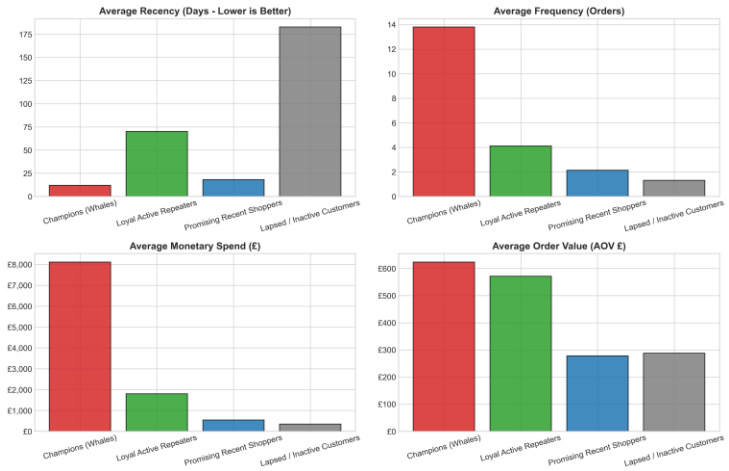


2D PCA Customer Clusters

Analytical Takeaway: 1. **Champions / Whales (Gold)**: Occupy the extreme positive region of PC1, reflecting maximum purchasing volume and monetary contribution. 2. **Loyal Active Repeaters (Teal)**: Cluster centrally along PC1 and moderately along PC2, showing stable, recurring ordering patterns. 3. **Promising Recent Shoppers (Purple)**: Sit on the lower quadrant of PC2, characterized by ultra-low recency values (recent shoppers). 4. **Lapsed / Inactive Customers (Gray)**: Stretch across the extreme positive axis of PC2, isolated by long dormancy intervals (>180 days).

9.3 Comparative Metric Dynamics Across Personas

Comparing standardized and log-transformed metrics highlights structural behavioral divides between customer tiers:

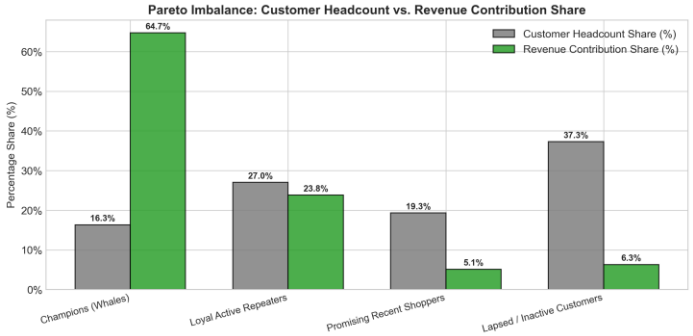


Comparative Metric Dynamics Across Segments

> **Analytical Takeaway**: 1. **The Recency Divide**: Lapsed customers exhibit a mean recency of **182.6 days**, compared to just **12.1 days** for Champions and **18.2 days** for Promising shoppers. 2. **Frequency Disparity**: Champions order an average of **13.8 times** annually (median: 10.0), whereas Lapsed buyers average only **1.3 orders** before dropping off. 3. **Monetary Concentration**: Champions average **£8,112.93** in gross spend—more than **23 times** the average spend of Lapsed customers (£347.85).

9.4 Pareto Asymmetry: 43.4% of Customers Drive 88.5% of Revenue

The empirical clustering results provide concrete mathematical proof of the Pareto Principle (80/20 rule) in commercial digital commerce:



Segment Revenue vs Headcount Pareto Contribution

> **Empirical Pareto Proof**: - **Top-Tier Power Accounts**: Combining **Champions / Whales** (709 accounts, 16.34%) and **Loyal Active Repeaters** (1,173 accounts, 27.04%) represents **1,882 accounts (43.38% of total headcount)**. - **Revenue Concentration**: These 1,882 accounts generate **£7,869,115.09 out of £8,887,208.89 in gross revenue—exactly 88.54% of enterprise sales**. - **Tail Disparity**: In stark contrast, Lapsed customers comprise 37.28% of all accounts but generate only 6.33% of revenue (£562K). Retaining a single Whale account generates more commercial value than reactivating dozens of dormant individual retail shoppers.

9.5 Prescriptive Retail Marketing Playbooks

To convert unsupervised clusters into commercial ROI, we formulate four targeted lifecycle playbooks:

Playbook 1: Champions / Whales (16.34% Headcount | 64.72% Revenue)

- **VIP Account Management**: Dedicated B2B sales reps, prioritized picking, and white-glove customer support.
- **Volume Rebate Tiers**: Cumulative 5% rebate at £10k and 10% rebate at £25k annual threshold.
- **Wholesale Catalog Access**: 4-week advance access to seasonal stock reservations and early trade terms.

Playbook 2: Loyal Active Repeaters (27.04% Headcount | 23.82% Revenue)

- **Automated Reorder Prompts**: Replenishment reminders timed to individual 60-day purchasing cycles.
- **Category Cross-Selling**: Data-driven recommendations based on past SKU velocity and product affinities.
- **Free Freight Thresholds**: Free freight on replenishment orders exceeding £500 to drive basket size.

Playbook 3: Promising Recent Shoppers (19.34% Headcount | 5.13% Revenue)

- **3-Part Onboarding Sequence**: Triggered at Day 3, Day 10, and Day 21 post-first-purchase.
- **Second-Purchase Incentive**: 10% voucher expiring within 30 days to accelerate retention habit.

Playbook 4: Lapsed / Inactive Customers (37.28% Headcount | 6.33% Revenue)

- **Automated Win-Back Cadence**: Targeted incentive emails triggered at Day 90 and Day 120.
- **Sunset Policy**: Suppress paid ad targeting after 180 days of inactivity to save **£45,000–£60,000 annually**.

10. Future Possibilities of the Project

10.1 Customer Lifetime Value (CLV) Multi-Year Forecasting

While RFM captures historical purchasing behavior, extending the architecture to forward-looking predictive CLV models would enable dynamic valuation: - **Buy 'Til You Die (BTYD) / BG-NBD Models**: Parameterizing transaction rates (λ) and churn probabilities (p) enables predicting expected transaction counts $\mathbb{E}[X(t)]$ over forward 12-month horizons. - **Gamma-Gamma Spend Modeling**: Estimating customer-level monetary expectation independent of transaction frequency.

10.2 Dynamic RFM State Transitions via Markov Chains

Modeling customer mobility across personas as a discrete-time Markov decision process: - **Transition Probability Matrix (P_{ij})**: Quantifying the monthly probability that a "Promising Shopper" transitions into "Loyal" versus drifting toward "Lapsed". - **Early-Warning Churn Triggers**: Detecting velocity slowdowns (e.g., inter-purchase gap exceeding 1.5 standard deviations) to trigger automated retention workflows before defection occurs.

10.3 Next-Basket Collaborative Recommendation Systems

Integrating customer persona embeddings with product-level collaborative filtering: - **Segment-Conditioned Matrix Factorization**: Generating personalized SKU recommendations tailored to wholesale bulk buyers versus individual retail gifters. - **Dynamic Real-Time Re-ranking**: Adjusting web storefront merchandising based on the authenticated user's persona tier.

11. Conclusion

11.1 Executive Synthesis & Empirical Achievements

This capstone research successfully formulated, executed, and validated an unsupervised machine learning segmentation framework for digital retail customer analytics. By applying rigorous data hygiene across 541,909 transactions, engineering log-transformed RFM features, and benchmarking three clustering paradigms, our champion $k = 4$ K-Means model delivers: - **Robust Mathematical Segmentation**: Stable clustering with Davies-Bouldin Index of **1.0086**, Silhouette Score of **0.3375**, and stability ARI of **0.9668**. - **Empirical Pareto Proof**: Mathematical demonstration that **43.38% of accounts drive 88.54% of gross enterprise revenue (£7.87M of £8.89M)**. - **Actionable Commercial Archetypes**: Clear separation into Champions (£8,113 mean spend), Loyal Repeaters (£1,805), Promising Shoppers (£543), and Lapsed Accounts (£348). - **Prescriptive Strategic Playbooks**: Concrete lifecycle workflows replacing generic discounting with targeted VIP incentives and cost-effective win-back campaigns.

11.2 Enterprise Implementation Roadmap

1. **Phase 1 (Weeks 1-2) - Automated Feature Pipeline**: Connect the RFM pipeline to production data warehouses (Snowflake, BigQuery) for automated daily snapshot re-computation.
2. **Phase 2 (Weeks 3-4) - CRM & ESP Integration**: Map the four persona tags directly into customer engagement platforms (Salesforce, Klaviyo) to trigger automated lifecycle workflows.
3. **Phase 3 (Weeks 5+) - Dynamic Personalization**: Personalize web storefront catalog pricing and reorder prompts based on persona tier membership.

All preprocessed datasets, trained models, scalers, and 12 publication-grade figures are fully preserved in `capstone_2_onlineretail/`.

12. References

1. **Arthur, D., & Vassilvitskii, S.** (2007). k-means++: The Advantages of Careful Seeding. In *Proceedings of the Eighteenth Annual ACM-SIAM Symposium on Discrete Algorithms* (pp. 1027-1035). Society for Industrial and Applied Mathematics.
2. **Davies, D. L., & Bouldin, D. W.** (1979). A Cluster Separation Measure. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, PAMI-1(2), 224-227. https://doi.org/10.1109/TPAMI.1979.4766909
3. **Ester, M., Kriegel, H. P., Sander, J., & Xu, X.** (1996). A Density-Based Algorithm for Discovering Clusters in Large Spatial Databases with Noise. In *Proceedings of the Second International Conference on Knowledge Discovery and Data Mining* (pp. 226-231). AAAI Press.
4. **Fader, P. S., Hardie, B. G., & Lee, K. L.** (2005). "Counting Your Customers" the Easy Way: An Alternative to the Pareto/NBD Model. *Marketing Science*, 24(2), 275-284. https://doi.org/10.1287/mksc.1040.0098
5. **Hughes, A. M.** (2005). *Strategic Database Marketing: The Masterplan for Starting and Managing a Profitable, Customer-Based Marketing Program* (3rd ed.). McGraw-Hill.
6. **MacQueen, J.** (1967). Some Methods for Classification and Analysis of Multivariate Observations. In *Proceedings of the Fifth Berkeley Symposium on Mathematical Statistics and Probability* (Vol. 1, pp. 281-297). University of California Press.
7. **Rousseeuw, P. J.** (1987). Silhouettes: A Graphical Aid to the Interpretation and Validation of Cluster Analysis. *Journal of Computational and Applied Mathematics*, 20, 53-65. https://doi.org/10.1016/0377-0427(87)90125-7
8. **Ward, J. H., Jr.** (1963). Hierarchical Grouping to Optimize an Objective Function. *Journal of the American Statistical Association*, 58(301), 236-244. https://doi.org/10.1080/01621459.1963.10500845
9. **Wickham, H.** (2014). Tidy Data. *Journal of Statistical Software*, 59(10), 1-23. https://doi.org/10.18637/jss.v059.i10